

Python in Healthcare

ETL Processing of Nested Data

May 11, 2017



Workshop: ETL Processing of Nested Data

ETL JSON Processing - Agenda

- Bio
- Problem Description
- Selected Python Techniques
- Transforming Nested Data
- Optimization Techniques
- Pitfalls
- Q & A

Hi, I'm Bijan!

- From Canada
- Started writing code at age 7
- B. Sc. in Software Engineering and Human Biology 2000
- Entered workforce as Software Developer since 1999
- MBA 2007
- Worked in FX/Derivative Trading, Gaming, D/R, and Network Apps
- Javascript, C/C++, Java, C#, perl, and ruby
- Worked on most modern O/S platforms
- Python developer since 2012
- Senior Developer with WebApps team @Clover Health since 2016

Source Code

Available here:

<https://github.com/CloverHealth/pycon2017>

Problem Description

Data collected via forms uses nested JSON in two shapes:

“Schema” describes Questions

```
{  
  "label": "Basic Info",  
  "slug": "basic_info",  
  "widget": "panel",  
  "nodeType": "section",  
  "children": [  
    {  
      "slug": "primary_phone",  
      "label": "Primary Phone",  
      "nodeType": "question",  
      "answerType": "text",  
      "widget": "phone_number_control",  
    },  
    ...  
  ]  
}
```

Basic Info

Member Information

First name

Clover AWW Test

Last name

Hunter

Member CPID

CP1150036

Date of birth

11/26/1950

Member Contact

Primary phone

415-555-5555

Problem Description

Data collected via forms uses nested JSON in two shapes:

“Submission” captures Answers

```
{
  "basic_info": {
    "member_info": {

      "vitals": {
        "name": "McTestface",
        "age": 42,
        "favorite_color": "blue",
        "bmi": 23
      }
    }
  }
}
```

Problem Description - Nesting

```
{
  "slug": "basic_info",
  ...
  "children": [
    {
      "slug": "member_info",
      ...
      "children": [
        {
          "slug": "contact_info",
          ...
          "children": [
            {
              "slug": "phone",
              ...
            },
            ...
          ]
        },
        ...
      ]
    },
    ...
  ]
}
```


Problem Description - Sample Output

| <i>schema_path (key)</i> | <i>value</i> |
|--|---------------------|
| basic_info.member_info.vitals. name | "McTestface" |
| basic_info.member_info.vitals. age | "42" |
| basic_info.member_info.vitals. favorite_color | "blue" |
| basic_info.member_info.vitals. bmi | "24" |
| . . . | . . . |
| other_form.new_section.info. bmi | "26" |

Problem Description - ETL Process

“ETL” = **Extract** + **Transform** + **Load**

- Extraction *Reading* source data
- Transform *Converting or computing* output data
- Load *Inserting* output data

Our ETL application will convert our data to a **flat** key/value table.

Why?

- Nested JSON is cumbersome to query in Postgres
- Different form schemas have different nested structures
- Single data format for easier “downstream” processing
 - e.g. metrics, reports, pivot views, etc.

Python Technologies Used

Core python

| | |
|----------------|--|
| Generators | Chaining flow of data Conserve memory |
| functools | Partial functions for encapsulation Caching |
| more_itertools | Batching / chunking |

11

Data Access

| | |
|--------------------|---|
| SQLAlchemy | Object Relational Mapper (ORM) |
| psycopg2 | Postgres database driver |
| testing.postgresql | Manage disposable/reuseable test database |

Python Generators

Generators are a powerful but simple technique for

- Concisely creating iterators
- Deferred evaluation
- Incremental evaluation
- Chaining operations

How?

- Use **yield** keyword to “generate” a value
- Returns an *iterable* object. Not a list!
- Evaluate at your own leisure
- Save memory.
- Easy to “chain” generators

Simple example: **squares()** generator function

Python Generators

Simple examples in `core_techniques.ipynb`

- `squares()` generator function
- Chaining generators (`fun_A()`, `fun_B()` and `fun_C()`)

Partial Function Application

Partial functions can be created in python using [`functools.partial\(\)`](#):

- “Freezes” some portion of a function’s args and/or kwargs
- Creates a new callable object
- Can still be called as a function
- Simplifies signature during execution

Back to examples in `core_techniques.ipynb`

Partial Function Application - Encapsulation

Encapsulation = “Construct that bundles data with methods”

You can do this in python using:

- | | |
|----------------------------------|--|
| • Classes | Data attributes in <code>__init__()</code> |
| • Closures / Inner methods | Variables from outer scope |
| • Partial (function) application | “Frozen” args/kwargs for a function |

These are all *valid* approaches.

Workshop code uses partial function application because:

- Quickly kwargs parameters from JSON config
- This *particular* application doesn't need class hierarchies
- Each of E, T and L only performs a *single operation*

Workshop Application - Core Design

- `app.etl.extractors` Read routines (extraction)
- `app.etl.transformers` JSON Transformation code
- `app.etl.loaders` Result insertion routines (loading)

- `app.processor` Main ETL flow in just 3 lines of code
 - Uses python generators to all 3 steps
 - Uses partial functions to encapsulate ETL parameters

See `core_techniques.ipynb` for simplified example

Transforming Nested Data

- Use a “mapping function” approach
 - Take iterable or generator
 - Apply functions to transform the data
 - Generate transformed results
- Known python built-in examples include:
 - [map\(\)](#) - applies a function to an iterable
 - [sorted\(\)](#) - applies a function to extract an ordering “key”
- For nested JSON structures, we need to:
 - Traverse the tree
 - Filter only required data from each tree node
 - Filter only required children

Transforming Nested Data

Workshop introduces `app.etl.transformers.map_nested()`

- INPUTS
 - python dict -- loaded JSON data
 - gen function to transform items at each node
 - gen function extract children to recurse into
- OUTPUTS:
 - transformed items (as a gen function)
- CORE ALGORITHM:
 - Pre-order DFS (depth-first search)
 - Current path (node slugs) are tracked using an array argument
 - Just 5 lines of code --- WAAAAAT?
 - Complexity resides in your INPUT generator functions

Transforming Nested Data - Quick App Demo

- Input Scenario: “demo”
 - 2 form schemas
 - 50 users
 - 50 submissions
- Processor setup: “naive-single”
 - Iterate on `Session.query()`
 - Call `Session.add()` for each response event
- Run:
 - `python main.py generate demo`
 - `python main.py process demo naive-single`

See `README.md` for detailed instructions

Transforming Nested Data - Performance

Even with generators, ETL processing of large amounts of data can incur performance problems

- ETL everything in memory?
 - ... but memory is limited
- ETL one row at a time?
 - ... will be very slow due to lots of I/O

Performance Optimization - Batching

Find a balance using **batches** (or “chunks”)

SQLAlchemy provides:

- [Query.yield_per\(\)](#) for extraction using database CURSORS
- [Session.bulk_save_objects\(\)](#) and [Session.bulk_insert_mappings\(\)](#) for load insertion

Performance Optimization - Batching

The “T” in ETL (the Transform) is not necessarily 1 to 1

- 1 row SELECTed will generally not lead to 1 row to INSERTed

In our workshop:

- One JSON structure will lead to *multiple* key-value pairs
- Different schemas will produce different number of key-value pairs

Do we need to write the “chunking” logic ourselves?

Performance Optimization - Batching

[more_itertools](#) library comes to the rescue!

- [chunked\(\)](#) method can take any input iterable
- Produces controlled chunks
- No complex logic needed

Loader can easily apply this to transformed output generator to reduce INSERTs

See `core_techniques.ipynb` for simplified example

[chunked\(\)](#) is just 1 line of code

... it uses a partial function!!

Performance Optimization - Caching

- For each form schema, the transformer computes a “path map”
 - `basic_info.member_info.name` -> text
 - `basic_info.member_info.age` -> number
- Optimize via caching by using [functools.lru_cache\(\)](#)
- But read the python docs carefully!
 - “the [...] arguments to the function must be hashable.”
- The SQLAlchemy “session” an opaque object.
 - Hashing would include internal state which changes
 - (loaded objects, cursors, etc.)

Performance Optimization - Caching

- No problem... use partial functions again!
- Create a custom function using [`functools.partial\(\)`](#)
 - Freeze the “session” argument (it’s an opaque reference)
 - `form_id` is then the only (hashable) parameter
- Wrap the resulting partial function with [`functools.lru_cache\(\)`](#)
- Client code just calls `get_node_path_map(form_id)`

Performance Pitfalls - SQLAlchemy ORM

SQLAlchemy ORM keeps code concise....

- Great for for general purpose in transactional apps
- e.g. Web sites, REST APIs, GraphQL
- Anything “CRUD”

... but gets tricky when dealing large datasets in OLAP

- e.g. Data Science, Analytics, Batch processing
- Default lazy loading leads to more SELECTs instead of joins (I/O)
- Transforming python objects to model instances (CPU)

Performance Pitfalls - SQLAlchemy ORM

For analytics, please consider any of these options:

- Explicitly specifying join load types if you must use the ORM
 - see `join_queries.ipynb` included
- SQLAlchemy Core to compose exact queries
 - Allows more precise queries
 - Be careful with readability!
- Using raw SQL
 - Best query precision (especially for advanced Postgres)
 - Often more readable
 - Via SQLAlchemy's [`Session.execute\(\)`](#)

Bonus! - Performance Analysis

This workshop code include performance test content!

- Timing and Memory Profiling
- Demonstrate tradeoffs between choices
- Repeatable due to “scenarios” template database

Run “jupyter notebook” to get detailed explanations on:

- **etl_analysis.ipynb** Performance tradeoffs
- **join_queries.ipynb** ORM join options for extraction

Please: Ask Bijan questions after the presentation!

Bonus! - Performance Analysis Tools

Profiling and Investigation

| | |
|------------------------|------------------|
| cProfile (core python) | Time profiling |
| memory_profiler | Memory profiling |

Presentation and Visualization

| | |
|------------|---|
| jupyter | Detailed write-up on analysis |
| gprof2dot | Visualize timing profiling (from cProfile) |
| matplotlib | Visualize memory profiling (from memory_profiler) |
| sqlparse | Pretty-printing raw SQL in logs (*) |

29

(*) NOTE: In healthcare, SQL debug poses risk of leaking protected information.

This workshop provides SQL debug logging, but all data is fake!

Bonus! - Performance Analysis Commands

You can analyze what is happening with the application:

- *Time* profiling

```
python -m cProfile -o timing.stats main.py process demo naive-single
```

```
bash visualize_pstats.sh timing.stats
```

- *Memory* profiling

```
mprof run python main.py process demo naive-single
```

```
mprof plot
```

- SQL debug logging

```
python main.py --debug-sql process demo naive-single > sql.log
```

```
less sql.log
```

Q & A

Credits

This workshop would not have been possible without prior work from these folks at Clover:

- James Bennett
- Joey Leingang
- David Flerlage
- Kathy Lass
- Diego Argueta
- Lavinia Karl
- Paul Minton

Thank you & contact information

Bijan Vakili

Senior Developer, Clover
Health

bijan.vakili@cloverhealth.com

Appendix

Testing Approach

Unit Tests

- Run all tests within DB transactions that are always rolled back
 - Use `Session.flush()` instead of `Session.commit()`
- Extend pytest CLI options to allow reuse a preset database
 - Uses the `base_dir` option in `testing.postgres`
 - Preset with schemas and no data
- `freezegun`: Artificially “Freeze” the system clock in unit tests

Performance Tests

- Reuse template database with preset datasets of any size
 - Uses the `copy_data_form` option in `testing.postgres`
 - Preset with schemas, forms and responses
 - No output events

Decision to use JSONB

- Faster reads for this workshop - no reparsing necessary
 - (but slower writes than JSON)
- No need to preserve semantically-insignificant whitespace between tokens
- No need to preserve key order
- Can support more Postgres operators and indexing

Pitfalls - SQLAlchemy, Postgres and timestamps

Be aware of how you use timestamps

- Always Include timezone
 - Stick with UTC wherever possible
- Watch out for time truncation by psycopg2 driver

Workshop Application - Usage

Execution: `python main.py ...`

- `generate` Create a template dataset (“scenario”)
- `process` Run the ETL transform using a scenario
- `psql` Connect to the resulting database and review the results
- `clean` Cleanup

Configuration: Edit `conf/...`

- `schemas/*.json` Sample JSON form schemas
- `perfddata.conf.json` Describe your scenarios
- `processors.conf.json` Tune your processor

See `README.md` for instructions

Transforming Nested Data

INPUT: Submissions (nested JSON responses)

OUTPUT: Response events (schema_path, value, tag)

app.etl.**transformers** module uses:

- Recursion to traverse nested subtrees
- Generators to yield events

Clover Platform - Common Elements

Infrastructure:

- Aptible for “Compliance as a Service”...
 - *PaaS Docker containers with healthcare compliance (HIPAA certified)*
 - Postgres DBs
 - SSL certificates
- ...on Amazon AWS
 - EC2, S3, Route 53, . . .

Critical Integrations

- Bug Capture/Monitoring - Sentry (raven)
- Monitoring - New Relic + Pager Duty (pygerduty)

Pitfalls - SQL logging

SQL debug logging is a great technique for finding root causes...

... but is discouraged in healthcare.

Sensitive member/patient information can be leaked in logs because of:

- INSERT parameters
- SELECT results
- “Data cleaning” tools exists... but customization effort is required
- This workshop’s code suppresses parameters and results in its logs
 - `app.log`
 - `app.util.sqldebug`